

C2_W4_Lab_3_using-TPU-strategy

February 5, 2025

1 TPU Strategy

In this ungraded lab you'll learn to set up the TPU Strategy. It is recommended you run this notebook in Colab by clicking the badge above. This will give you access to a TPU as mentioned in the walkthrough video. Make sure you set your runtime to TPU.

1.1 Imports

```
[1]: import os
import random
try:
    # %tensorflow_version only exists in Colab.
    %tensorflow_version 2.x
except Exception:
    pass

import tensorflow as tf
print("TensorFlow version " + tf.__version__)
AUTO = tf.data.experimental.AUTOTUNE
```

TensorFlow version 2.3.1

1.2 Set up TPUs and initialize TPU Strategy

Ensure to change the runtime type to TPU in Runtime -> Change runtime type -> TPU

```
[2]: # Detect hardware
try:
    tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver(tpu_address) # TPU_
    ↪ detection
    tf.config.experimental_connect_to_cluster(tpu)
    tf.tpu.experimental.initialize_tpu_system(tpu)
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
    # Going back and forth between TPU and host is expensive.
    # Better to run 128 batches on the TPU before reporting back.
```

```

print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
print("Number of accelerators: ", strategy.num_replicas_in_sync)
except ValueError:
    print('TPU failed to initialize.')

```

```

↳ -----

KeyError                                Traceback (most recent call↳
↳ last)

<ipython-input-2-58a4c002f41d> in <module>
      1 # Detect hardware
      2 try:
----> 3     tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
      4     tpu = tf.distribute.cluster_resolver.
↳ TPUClusterResolver(tpu_address) # TPU detection
      5     tf.config.experimental_connect_to_cluster(tpu)

/opt/conda/lib/python3.7/os.py in __getitem__(self, key)
    677         except KeyError:
    678             # raise KeyError with the original key value
--> 679             raise KeyError(key) from None
    680         return self.decodevalue(value)
    681

KeyError: 'COLAB_TPU_ADDR'

```

1.3 Download the Data from Google Cloud Storage

```

[ ]: SIZE = 224 #@param ["192", "224", "331", "512"] {type:"raw"}
    IMAGE_SIZE = [SIZE, SIZE]

```

```

[ ]: GCS_PATTERN = 'gs://flowers-public/tfrecords-jpeg-{}x{}/*.tfrec'.
    ↳ format(IMAGE_SIZE[0], IMAGE_SIZE[1])

    BATCH_SIZE = 128 # On TPU in Keras, this is the per-core batch size. The
    ↳ global batch size is 8x this.

    VALIDATION_SPLIT = 0.2
    CLASSES = ['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips'] # do not
    ↳ change, maps to the labels in the data (folder names)

```

```

# splitting data files between training and validation
filenames = tf.io.gfile.glob(GCS_PATTERN)
random.shuffle(filenames)

split = int(len(filenames) * VALIDATION_SPLIT)
training_filenames = filenames[split:]
validation_filenames = filenames[:split]
print("Pattern matches {} data files. Splitting dataset into {} training files,
      {} and {} validation files".format(len(filenames), len(training_filenames),
      len(validation_filenames)))

validation_steps = int(3670 // len(filenames) * len(validation_filenames)) // BATCH_SIZE
steps_per_epoch = int(3670 // len(filenames) * len(training_filenames)) // BATCH_SIZE
print("With a batch size of {}, there will be {} batches per training epoch and {} batch(es) per validation run.".format(BATCH_SIZE, steps_per_epoch, validation_steps))

```

1.4 Create a dataset from the files

- `load_dataset` takes the filenames and turns them into a `tf.data.Dataset`
- `read_tfrecord` parses out a tf record into the image, class and a one-hot-encoded version of the class
- Batch the data into training and validation sets with helper functions

```

[ ]: def read_tfrecord(example):
    features = {
        "image": tf.io.FixedLenFeature([], tf.string), # tf.string means
        bytestring
        "class": tf.io.FixedLenFeature([], tf.int64), # shape [] means scalar
        "one_hot_class": tf.io.VarLenFeature(tf.float32),
    }
    example = tf.io.parse_single_example(example, features)
    image = example['image']
    class_label = example['class']
    image = tf.image.decode_jpeg(image, channels=3)
    image = tf.image.resize(image, [224, 224])
    image = tf.cast(image, tf.float32) / 255.0 # convert image to floats in
    [0, 1] range
    class_label = tf.cast(class_label, tf.int32)
    return image, class_label

def load_dataset(filenames):

```

```

# read from TFRecords. For optimal performance, use "interleave(tf.data.
↪TFRecordDataset, ...)"
# to read from multiple TFRecord files at once and set the option
↪experimental_deterministic = False
# to allow order-altering optimizations.

option_no_order = tf.data.Options()
option_no_order.experimental_deterministic = False

dataset = tf.data.Dataset.from_tensor_slices(filenamees)
dataset = dataset.with_options(option_no_order)
dataset = dataset.interleave(tf.data.TFRecordDataset, cycle_length=16,
↪num_parallel_calls=AUTO) # faster
dataset = dataset.map(read_tfrecord, num_parallel_calls=AUTO)
return dataset

def get_batched_dataset(filenamees):
    dataset = load_dataset(filenamees)
    dataset = dataset.shuffle(2048)
    dataset = dataset.batch(BATCH_SIZE, drop_remainder=False) # drop_remainder
↪will be needed on TPU
    dataset = dataset.prefetch(AUTO) # prefetch next batch while training
↪(autotune prefetch buffer size)
    return dataset

def get_training_dataset():
    dataset = get_batched_dataset(training_filenames)
    dataset = strategy.experimental_distribute_dataset(dataset)
    return dataset

def get_validation_dataset():
    dataset = get_batched_dataset(validation_filenames)
    dataset = strategy.experimental_distribute_dataset(dataset)
    return dataset

```

1.5 Define the Model and training parameters

```

[ ]: class MyModel(tf.keras.Model):
    def __init__(self, classes):
        super(MyModel, self).__init__()
        self._conv1a = tf.keras.layers.Conv2D(kernel_size=3, filters=16,
↪padding='same', activation='relu')
        self._conv1b = tf.keras.layers.Conv2D(kernel_size=3, filters=30,
↪padding='same', activation='relu')
        self._maxpool1 = tf.keras.layers.MaxPooling2D(pool_size=2)

```

```

        self._conv2a = tf.keras.layers.Conv2D(kernel_size=3, filters=60,
        ↳padding='same', activation='relu')
        self._maxpool2 = tf.keras.layers.MaxPooling2D(pool_size=2)

        self._conv3a = tf.keras.layers.Conv2D(kernel_size=3, filters=90,
        ↳padding='same', activation='relu')
        self._maxpool3 = tf.keras.layers.MaxPooling2D(pool_size=2)

        self._conv4a = tf.keras.layers.Conv2D(kernel_size=3, filters=110,
        ↳padding='same', activation='relu')
        self._maxpool4 = tf.keras.layers.MaxPooling2D(pool_size=2)

        self._conv5a = tf.keras.layers.Conv2D(kernel_size=3, filters=130,
        ↳padding='same', activation='relu')
        self._conv5b = tf.keras.layers.Conv2D(kernel_size=3, filters=40,
        ↳padding='same', activation='relu')

        self._pooling = tf.keras.layers.GlobalAveragePooling2D()
        self._classifier = tf.keras.layers.Dense(classes, activation='softmax')

    def call(self, inputs):
        x = self._conv1a(inputs)
        x = self._conv1b(x)
        x = self._maxpool1(x)

        x = self._conv2a(x)
        x = self._maxpool2(x)

        x = self._conv3a(x)
        x = self._maxpool3(x)

        x = self._conv4a(x)
        x = self._maxpool4(x)

        x = self._conv5a(x)
        x = self._conv5b(x)

        x = self._pooling(x)
        x = self._classifier(x)
        return x

```

```

[ ]: with strategy.scope():
    model = MyModel(classes=len(CLASSES))
    # Set reduction to `none` so we can do the reduction afterwards and divide by
    # global batch size.
    loss_object = tf.keras.losses.SparseCategoricalCrossentropy(

```

```

reduction=tf.keras.losses.Reduction.NONE)

def compute_loss(labels, predictions):
    per_example_loss = loss_object(labels, predictions)
    return tf.nn.compute_average_loss(per_example_loss,
↪global_batch_size=BATCH_SIZE * strategy.num_replicas_in_sync)

test_loss = tf.keras.metrics.Mean(name='test_loss')

train_accuracy = tf.keras.metrics.SparseCategoricalAccuracy(
    name='train_accuracy')
test_accuracy = tf.keras.metrics.SparseCategoricalAccuracy(
    name='test_accuracy')

optimizer = tf.keras.optimizers.Adam()

@tf.function
def distributed_train_step(dataset_inputs):
    per_replica_losses = strategy.run(train_step,args=(dataset_inputs,))
    print(per_replica_losses)
    return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses,
                           axis=None)

@tf.function
def distributed_test_step(dataset_inputs):
    strategy.run(test_step, args=(dataset_inputs,))

def train_step(inputs):
    images, labels = inputs

    with tf.GradientTape() as tape:
        predictions = model(images)
        loss = compute_loss(labels, predictions)

    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))

    train_accuracy.update_state(labels, predictions)

    return loss

def test_step(inputs):
    images, labels = inputs

    predictions = model(images)
    loss = loss_object(labels, predictions)

```

```
test_loss.update_state(loss)
test_accuracy.update_state(labels, predictions)
```

```
[ ]: EPOCHS = 40
with strategy.scope():
    for epoch in range(EPOCHS):
        # TRAINING LOOP
        total_loss = 0.0
        num_batches = 0
        for x in get_training_dataset():
            total_loss += distributed_train_step(x)
            num_batches += 1
        train_loss = total_loss / num_batches

        # TESTING LOOP
        for x in get_validation_dataset():
            distributed_test_step(x)

        template = ("Epoch {}, Loss: {:.2f}, Accuracy: {:.2f}, Test Loss: {:.2f}, "
                    "Test Accuracy: {:.2f}")
        print (template.format(epoch+1, train_loss,
                                train_accuracy.result()*100, test_loss.result() /
→strategy.num_replicas_in_sync,
                                test_accuracy.result()*100))

        test_loss.reset_states()
        train_accuracy.reset_states()
        test_accuracy.reset_states()
```

1.6 Predictions

```
[ ]: #@title display utilities [RUN ME]
import matplotlib.pyplot as plt

def dataset_to_numpy_util(dataset, N):
    dataset = dataset.batch(N)

    if tf.executing_eagerly():
        # In eager mode, iterate in the Dataset directly.
        for images, labels in dataset:
            numpy_images = images.numpy()
            numpy_labels = labels.numpy()
            break;

    else: # In non-eager mode, must get the TF note that
```

```

        # yields the next item and run it in a tf.Session.
        get_next_item = dataset.make_one_shot_iterator().get_next()
        with tf.Session() as ses:
            numpy_images, numpy_labels = ses.run(get_next_item)

    return numpy_images, numpy_labels

def title_from_label_and_target(label, correct_label):
    label = np.argmax(label, axis=-1) # one-hot to class number
    # correct_label = np.argmax(correct_label, axis=-1) # one-hot to class number
    correct = (label == correct_label)
    return "{} [{}{}{}].format(CLASSES[label], str(correct), ', should be ' if
    ↪not correct else '',
                                CLASSES[correct_label] if not correct else ''),
    ↪correct

def display_one_flower(image, title, subplot, red=False):
    plt.subplot(subplot)
    plt.axis('off')
    plt.imshow(image)
    plt.title(title, fontsize=16, color='red' if red else 'black')
    return subplot+1

def display_9_images_from_dataset(dataset):
    subplot=331
    plt.figure(figsize=(13,13))
    images, labels = dataset_to_numpy_util(dataset, 9)
    for i, image in enumerate(images):
        title = CLASSES[np.argmax(labels[i], axis=-1)]
        subplot = display_one_flower(image, title, subplot)
        if i >= 8:
            break;

    plt.tight_layout()
    plt.subplots_adjust(wspace=0.1, hspace=0.1)
    plt.show()

def display_9_images_with_predictions(images, predictions, labels):
    subplot=331
    plt.figure(figsize=(13,13))
    for i, image in enumerate(images):
        title, correct = title_from_label_and_target(predictions[i], labels[i])
        subplot = display_one_flower(image, title, subplot, not correct)
        if i >= 8:
            break;

    plt.tight_layout()

```



```
plt.subplots_adjust(wspace=0.1, hspace=0.1)
plt.show()
```

```
def display_training_curves(training, validation, title, subplot):
    if subplot%10==1: # set up the subplots on the first call
        plt.subplots(figsize=(10,10), facecolor='#F0F0F0')
        plt.tight_layout()
    ax = plt.subplot(subplot)
    ax.set_facecolor('#F8F8F8')
    ax.plot(training)
    ax.plot(validation)
    ax.set_title('model ' + title)
    ax.set_ylabel(title)
    ax.set_xlabel('epoch')
    ax.legend(['train', 'valid.'])
```

```
[ ]: inference_model = model
```

```
[ ]: some_flowers, some_labels = □
    ↪ dataset_to_numpy_util(load_dataset(validation_filenames), 8*20)
```

```
[ ]: import numpy as np
    # randomize the input so that you can execute multiple times to change results
    permutation = np.random.permutation(8*20)
    some_flowers, some_labels = (some_flowers[permutation], □
    ↪ some_labels[permutation])

    predictions = inference_model(some_flowers)

    print(np.array(CLASSES)[np.argmax(predictions, axis=-1)].tolist())

    display_9_images_with_predictions(some_flowers, predictions, some_labels)
```